

Association Rule Mining and Twitter

Twitter API Setup

To access a Twitter API you will need to set up an account and receive a consumerKey, the consumerSecret, the access_Token, and the access_Secret. A popular library and API: "twitterR".

```
knitr::opts_chunk$set(echo = TRUE, warning=FALSE, results = 'show', include=TRUE, messages=FALSE)
```

```
##### Twitter in R
# Consumer API keys
# Access token & access token secret

## I have created a text file that contains the
## consumerKey, the consumerSecret, the access_Token, and the access_Secret
## They are comma seperated.
# Insert your consumerKey and consumerSecret below

consumerKey='SiMslBfTdWEimvLweRDTTrZVH'
consumerSecret='FoPYqK3uwpzutwE6G1RmQvPbRJ8RChFSLfIlgRAcFHjymKDzHh'
access_Token='1084502204038479872-v2czQaDlMt9ikoLnxhiQYk8Yb3f0RT'
access_Secret='U9ktzvd5rEwcK13mttsgwAujS0VxNPtJstxXcEE5znnid'
```

Once you have your keys, you can set up the API.

```
requestURL='https://api.twitter.com/oauth/request_token'
accessURL='https://api.twitter.com/oauth/access_token'
authURL='https://api.twitter.com/oauth/authorize'

### NOTES: rtweet is another excellent option
## https://mkearney.github.io/blog/2017/06/01/intro-to-rtweet/
### https://rtweet.info/

### Install the needed packages...
#install.packages("twitterR")
#install.packages("ROAuth")
# install.packages("rtweet")
library(arules)

## Loading required package: Matrix
##
## Attaching package: 'arules'
## The following objects are masked from 'package:base':
##
##      abbreviate, write

library(rtweet)
library(twitterR)

##
## Attaching package: 'twitterR'
```

```

## The following object is masked from 'package:rtweet':
##
##      lookup_statuses
library(ROAuth)
library(jsonlite)

##
## Attaching package: 'jsonlite'
## The following object is masked from 'package:rtweet':
##
##      flatten
#install.packages("streamR")
#library(streamR)
#install.packages("rjson")
library(rjson)

##
## Attaching package: 'rjson'
## The following objects are masked from 'package:jsonlite':
##
##      fromJSON, toJSON
#install.packages("tokenizers")
library(tokenizers)
library(tidyverse)

## -- Attaching packages ----- tidyverse_
## v ggplot2 3.2.1      v purrr   0.3.2
## v tibble  2.1.3      v dplyr   0.8.1
## v tidyr   0.8.3      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0

## -- Conflicts ----- tidyverse_c
## x tidyr::expand()    masks Matrix::expand()
## x dplyr::filter()    masks stats::filter()
## x purrr::flatten()   masks jsonlite::flatten(), rtweet::flatten()
## x rjson::fromJSON()  masks jsonlite::fromJSON()
## x dplyr::id()        masks twitterR::id()
## x dplyr::lag()       masks stats::lag()
## x dplyr::location()  masks twitterR::location()
## x dplyr::recode()    masks arules::recode()
## x rjson::toJSON()    masks jsonlite::toJSON()
library(plyr)

## -----
## You have loaded plyr after dplyr - this is likely to cause problems.
## If you need functions from both plyr and dplyr, please load plyr first, then dplyr:
## library(plyr); library(dplyr)
## -----
##
## Attaching package: 'plyr'

```

```
## The following objects are masked from 'package:dplyr':
##
##   arrange, count, desc, failwith, id, mutate, rename, summarise,
##   summarize
## The following object is masked from 'package:purrr':
##
##   compact
## The following object is masked from 'package:twitter':
##
##   id
```

```
library(dplyr)
library(ggplot2)
#install.packages("syuzhet") ## sentiment analysis
#library(syuzhet)
library(stringr)
#install.packages("arulesViz")
library(arulesViz)
```

```
## Loading required package: grid
library(semPlot)
```

Collecting Tweets

Next we will set up the API and search for a particular hash tag. We will store the tweets with the designated hash in a csv file for safe keeping. Here, we choose “Trump” in hopes to get a 100 tweets easily.

```
##### Using twittR #####
setup_twitter_oauth(consumerKey,consumerSecret,access_Token,access_Secret)
```

```
## [1] "Using direct authentication"
```

```
Search<-twitterR::searchTwitter("nfl",n=90,since="2019-10-14")
Search_DF <- twListToDF(Search)
TransactionTweetsFile = "Choc.csv"
#Search_DF$text[1]
```

```
## Start the file
Trans <- file(TransactionTweetsFile)
## Tokenize to words
Tokens<-tokenizers::tokenize_words(Search_DF$text[1],stopwords = stopwords::stopwords("en"),
  lowercase = TRUE, strip_punct = TRUE, strip_numeric = TRUE,simplify = TRUE)
## Write squished tokens
cat(unlist(str_squish(Tokens)), "\n", file=Trans, sep=",")
close(Trans)
```

```
## Append remaining lists of tokens into file
## Recall - a list of tokens is the set of words from a Tweet
Trans <- file(TransactionTweetsFile, open = "a")
for(i in 2:nrow(Search_DF)){
  Tokens<-tokenize_words(Search_DF$text[i],stopwords = stopwords::stopwords("en"),
    lowercase = TRUE, strip_punct = TRUE, simplify = TRUE)
```

```

    cat(unlist(str_squish(Tokens)), "\n", file=Trans, sep=",")
}
close(Trans)

```

Tweets as Transactions

In this section we will read in the tweets stored in the CSV file using the (Association Rule Mining) ARM library. Each tweet will be considered a basket of words. We can use ARM to determine associations of words in tweets.

```

##### Read in the tweet transactions
TweetTrans <- read.transactions(TransactionTweetsFile,
                                rm.duplicates = FALSE,
                                format = "basket",
                                sep=",",
                                ## cols =
                                )

#inspect(TweetTrans)
## See the words that occur the most
Sample_Trans <- sample(TweetTrans, 20)
#summary(Sample_Trans)

## Read the transactions data into a dataframe
TweetDF <- read.csv(TransactionTweetsFile, header = FALSE, sep = ",")
head(TweetDF)

```

```

##      V1      V2      V3      V4      V5      V6
## 1   rt      nfl    welcome      ramsnfl jalenramsey https
## 2   rt      lbnfl  reminder      bumpnrungilm0re      best  cb
## 3 final      decision retirement robgronkowski<U+2069>      https t.co
## 4   rt mreeseeagles officiating      talked      part  nfl
## 5   rt bravovictor03      nfl      waits      9 months
## 6   rt sharplinesdfs      hour      left      till  nfl
##      V7      V8      V9      V10     V11     V12     V13 V14     V15
## 1      t.co yuk63w2yyd
## 2      nfl      https      t.co coh1ytazjb
## 3 dftxlxu5ja
## 4  football  yellow  flags  impacting games  like never
## 5      amp      6    games  regular season saints  get  5    1
## 6      lock      time research      let    us  work  join  us just
##  V16     V17     V18 V19
## 1
## 2
## 3
## 4
## 5 now decide suspend
## 6 25  month      get

```

```
#(str(TweetDF))
```

Cleaning the text data

Note that cleaning the text data is very important in text mining applications. Tweets are especially “messy”. We will remove “rt”, “http”, etc and any other strings of no importance.

```
## Convert all columns to char
```

```
TweetDF<-TweetDF %>%  
  mutate_all(as.character)  
(str(TweetDF))
```

```
## 'data.frame':   95 obs. of  19 variables:  
## $ V1 : chr  "rt" "rt" "final" "rt" ...  
## $ V2 : chr  "nfl" "lbnfl" "decision" "mreeseeagles" ...  
## $ V3 : chr  "welcome" "reminder" "retirement" "officiating" ...  
## $ V4 : chr  "ramsnfl" "bumpnrungilm0re" "robgronkowski<U+2069>" "talked" ...  
## $ V5 : chr  "jalenramsey" "best" "https" "part" ...  
## $ V6 : chr  "https" "cb" "t.co" "nfl" ...  
## $ V7 : chr  "t.co" "nfl" "dftxlxu5ja" "football" ...  
## $ V8 : chr  "yuk63w2yyd" "https" "" "yellow" ...  
## $ V9 : chr  "" "t.co" "" "flags" ...  
## $ V10: chr  "" "coh1ytazjb" "" "impacting" ...  
## $ V11: chr  "" "" "" "games" ...  
## $ V12: chr  "" "" "" "like" ...  
## $ V13: chr  "" "" "" "never" ...  
## $ V14: chr  "" "" "" "" ...  
## $ V15: chr  "" "" "" "" ...  
## $ V16: chr  "" "" "" "" ...  
## $ V17: chr  "" "" "" "" ...  
## $ V18: chr  "" "" "" "" ...  
## $ V19: chr  "" "" "" "" ...
```

```
## NULL
```

```
# We can now remove certain words
```

```
TweetDF[TweetDF == "t.co"] <- ""  
TweetDF[TweetDF == "rt"] <- ""  
TweetDF[TweetDF == "http"] <- ""  
TweetDF[TweetDF == "https"] <- ""  
TweetDF[TweetDF == "sxrgihoe"] <- ""
```

```
## Clean with grepl - every row in each column
```

```
MyDF<-NULL
```

```
for (i in 1:ncol(TweetDF)){  
  MyList=c() # each list is a column of logicals ...  
  MyList=c(MyList,grepl("[:digit:]", TweetDF[[i]]))  
  MyDF<-cbind(MyDF,MyList) ## create a logical DF  
  ## TRUE is when a cell has a word that contains digits  
}
```

```
## For all TRUE, replace with blank
```

```
TweetDF[MyDF] <- ""  
(head(TweetDF,10))
```

```
##          V1          V2          V3          V4          V5          V6
```

```

## 1          nfl      welcome ramsnfl jalenramsey
## 2          lbnfl      reminder          best          cb
## 3  final          decision  retirement
## 4          mreeseeagles officiating  talked          part          nfl
## 5          nfl      waits          months
## 6          sharplinesdfs          hour  left          till          nfl
## 7          simmons_szn          almost  time          year          nfl
## 8          seahawks          vote          week's          ground
## 9  really          kills          team  really
## 10         nfl_memes          nfl      refs          like realdockery
##          V7      V8          V9          V10  V11  V12  V13 V14  V15 V16
## 1
## 2      nfl
## 3
## 4  football  yellow      flags  impacting  games  like never
## 5      amp          games      regular season saints  get          now
## 6      lock      time research          let      us  work  join  us just
## 7      nba college football basketball
## 8      player      week
## 9          wasn't      refs          maybe couple  plays  coul
## 10
##          V17      V18 V19
## 1
## 2
## 3
## 4
## 5  decide suspend
## 6  month      get
## 7
## 8
## 9
## 10

```

```

# Now we save the dataframe using the write table command
write.table(TweetDF, file = "UpdatedChocolate.csv", col.names = FALSE,
            row.names = FALSE, sep = ",")
TweetTrans <- read.transactions("UpdatedChocolate.csv", sep = ",",
                                format("basket"), rm.duplicates = TRUE)

```

```

## distribution of transactions with duplicates:
## items
## 1 2 3
## 10 3 1

```

```

#inspect(TweetTrans)

```

ARM

Next we will apply the apriori algorithm to find the associations including computing the support, confidence and lift. Read more on the arules library to tweak / tune the following code to achieve desired results.

```

# So that you do not have an enormous amount of rules, you can thresholds for
# support, confidence and lift ... also minlength for the rules.
TweetTrans_rules = arules::apriori(TweetTrans,

```

```

parameter = list(support=0.05, confidence=.65, minlen=3))

## Apriori
##
## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
##          0.65    0.1    1 none FALSE                TRUE      5    0.05    3
## maxlen target  ext
##          10  rules FALSE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
##      0.1 TRUE TRUE  FALSE TRUE    2    TRUE
##
## Absolute minimum support count: 4
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[522 item(s), 95 transaction(s)] done [0.00s].
## sorting and recoding items ... [15 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 done [0.00s].
## writing ... [3 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].
inspect(head(TweetTrans_rules, 10))

##      lhs                      rhs          support    confidence
## [1] {rapsheet,tompelissero} => {nfl}      0.05263158 1.0000000
## [2] {nfl,tompelissero}      => {rapsheet}  0.05263158 1.0000000
## [3] {nfl,rapsheet}          => {tompelissero} 0.05263158 0.7142857
##      lift      count
## [1]  1.397059  5
## [2] 13.571429  5
## [3] 13.571429  5

## sorted
SortedRules_conf <- sort(TweetTrans_rules, by="confidence", decreasing=TRUE)
inspect(head(SortedRules_conf, 10))

##      lhs                      rhs          support    confidence
## [1] {rapsheet,tompelissero} => {nfl}      0.05263158 1.0000000
## [2] {nfl,tompelissero}      => {rapsheet}  0.05263158 1.0000000
## [3] {nfl,rapsheet}          => {tompelissero} 0.05263158 0.7142857
##      lift      count
## [1]  1.397059  5
## [2] 13.571429  5
## [3] 13.571429  5

SortedRules_sup <- sort(TweetTrans_rules, by="support", decreasing=TRUE)
inspect(head(SortedRules_sup, 10))

##      lhs                      rhs          support    confidence
## [1] {rapsheet,tompelissero} => {nfl}      0.05263158 1.0000000
## [2] {nfl,tompelissero}      => {rapsheet}  0.05263158 1.0000000
## [3] {nfl,rapsheet}          => {tompelissero} 0.05263158 0.7142857
##      lift      count

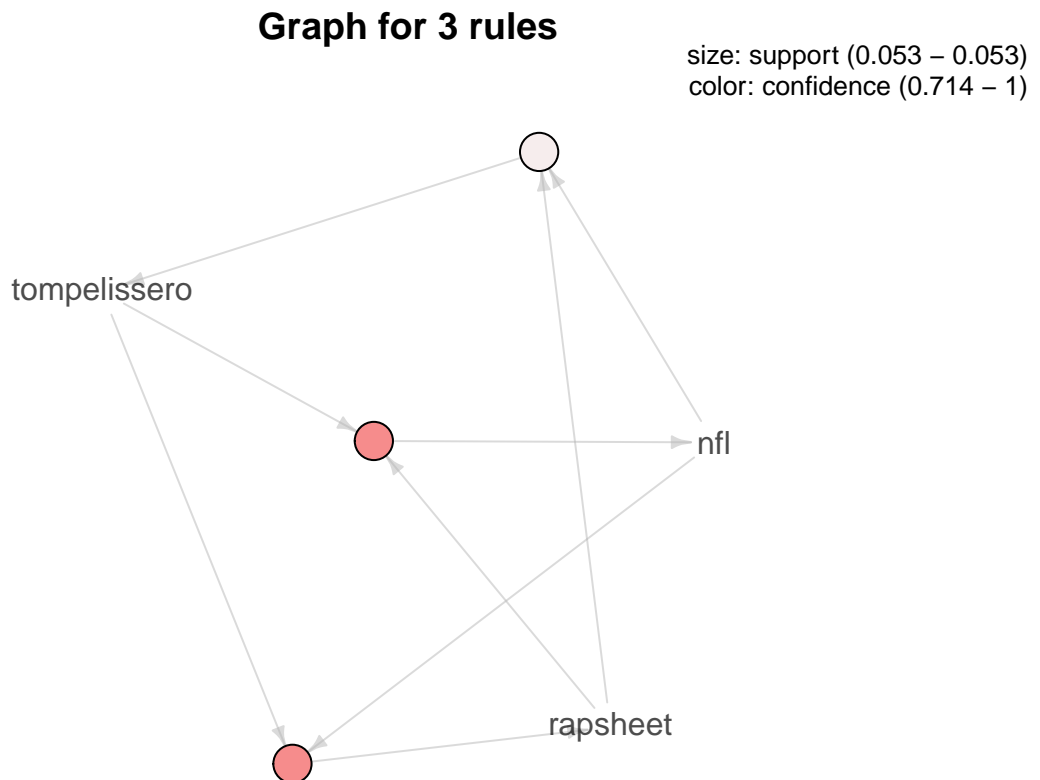
```

```
## [1] 1.397059 5
## [2] 13.571429 5
## [3] 13.571429 5
```

Displaying Results

The results will be displayed as an interactive graph.

```
plot (head(SortedRules_sup,n=10),method="graph",shading="confidence")
```



```
plot (head(SortedRules_conf, n=10),method="graph",shading="confidence")
```


Graph for 3 rules

size: support (0.053 – 0.053)
color: confidence (0.714 – 1)

