

Assignment for Week 9 / Modules 5 and 6 (Merged and consolidated)

In this assignment, you'll be asked to predict reading scores using the ELS data.

Scatterplots.

1. Choose a dependent variable and plot that dependent variable as a function of a continuous independent variable.

Use `ggplot` ☺

2. If you need to, adjust the scale of either the dependent variable or independent variable to make more sense.

Yes – try this and show the results (even if it makes the plot “worse”).

3. Add a line of best fit to your graphic.

Show both the loess and lm lines. Hint: Use `geom smooth`.

4. Clean up the labeling and description of the graphic so that it could be presented to a broader audience.

Make it pretty. Hint: Investigate the resources for `ggplot` and find which parameters can be altered for aesthetics.

5. In your .Rmd file, include a few sentences describing what the graphic above shows.

Also, clearly explain the difference between the loess and the lm estimates. Explain which would be “better” for prediction purposes and why.

Regression.

6. Create a regression that predicts reading scores as a function of SES **plus one other covariate**, using the training data.

First load the training data and proceed with “learning” a regression model. This is synonymous with learning the parameters to the linear equation, such that the linear equation will “best” model the data. Hint: use the `lm(...)` function; you will need to clearly identify the independent and dependent variable(s)

Hint: I encourage you to plot the results so that you can visually confirm the result.

7. Report the RMSE from a validation of your model using the testing data.

First you will need to load the testing data. Then apply your learned model (from #1) to the test data. The prediction will likely not be perfect, and thus there will be some error ... that's what we are computing here !! Hint: use the `modelr::rmse(...)` function. This function will apply the regression model and measure the error!

Hint: (As you have done previously) You may wish to store intermediate computational results of the error back into the data frame using `mutate`.

Hint: I encourage you to plot the results so that you can visually confirm the result.